

# EMBS Autumn Assessment

Y6385133

Autumn 2012

## 1 Modelling

I chose to model the solution using Mote Runner SDK. The main reasons against using another solution were:

### Differences in implementation details

There is no guarantee that applications such as ptollemey would exhibit the same characteristics as a mote

### Time constraints

I expected the majority of my time would be taken up with debugging the mote SDK

## 2 Implementation

My original implementation observed a channel until it had determined the period, and then analysed the next sink. This worked, but in the worst case situation where the  $Sink_A$  has  $n = 1500$  and  $t = 10$  the source could block for 35s – over half the time of the entire simulation – before it analyses other sinks. See figure 1.

To fix this I added a timeout that restricts the time the source can spend trying to analyse a channel when there are other periods that need estimating.

The general case for determining a sink's period  $t$  involves observing two of the sequence numbers  $(n_1, n_2)$  emitted by a sink, recording the times they were observed  $(r_1, r_2)$  and then calculating:

$$t = \frac{r_2 - r_1}{n_1 - n_2}$$

This allows the period to be calculated without observing consecutive sequence numbers, however the two sequence numbers must be observed within the same synchronisation phase. In the case where  $n = 1$  this is impossible, so I added a special case that calculates

$$t = \frac{r_2 - r_1}{12}$$




Figure 1: Without a timeout the source can block for too long and miss other phases

However, this can be problematic if the  $n_2$  is not in the cycle after  $n_1$ 's. One countermeasure is to ignore periods  $t > 1600$  (+100ms to allow for clock drift)

When  $n_2 > n_1$  we cannot reliably calculate the period as the second sequence number may not be the sink's  $n_{max}$ , thus giving a larger  $t$ , which could cause us to transmit outside of the reception period. In this instance  $n_1$  is replaced with  $n_2$ 's values and  $n_2$  is reset.

There is a hard limit of 1600ms to prevent obviously erroneous periods.

### 3 Clock Drift

This is inevitable in systems without a centralised clock, especially when wireless communication is involved. To help combat drift the source always performs calculations using times events were triggered at, rather than 'current' time. After each broadcast the source schedules a new broadcast in  $t * (11 + n)$  ticks and a resync in  $(t * (10 + n)) - x$  ticks. The resync is used to verify that the estimated reception period is accurate. Unfortunately, this is occasionally missed due to competing transmission slots.

### 4 Power Saving

My implementation saves power by scaling down the transmission power based on the signal strength of received packets and turning off LEDs (when not debugging). Received power strength is between 0 – 255 and transmission power 0 – 63. Calculating min power needed to send to  $sink_x$ :

$$tx\_power_x = \frac{rx\_power_x}{255} \times 63$$




Figure 2: Bug in power consumption measurements

However using this exact value will mean a different min power is required as noise/distance from sink changes. Instead I transmit at double the minimum required power.

$$tx\_power_x = \max(63, \frac{rx\_power_x}{255} \times 63)$$

The above operation can be approximated using multiplications of powers of 2, aka bit shifts.

$$\begin{aligned} tx\_power_x &= rx\_power_x \times 2^{-8} \times 2^6 \\ &= rx\_power_x \times 2^{-2} \end{aligned}$$

## 5 Verification

The SDK's netview tool was used to confirm reduction in TX power, however it appears there's a bug that reports an extreme increase in current when signal strengths other than max are used (figure 3)

Ignoring erroneous values yields figure ??.

I also used a script by another student (not included) to generate sinks with different parameters, and automated it to test for multiple scenarios:

These values were all observed in the simulator so performance on hardware may be different.




Figure 3: Consumption sans bugged values. Note that PSM uses on average less power than max power

$Sink_A$				$Sink_B$				$Sink_C$			
$n$	$t(ms)$	$A$	$R$	$n$	$t(ms)$	$A$	$R$	$n$	$t(ms)$	$A$	$R$
10	500	6	0	4	700	6	0	5	1500	3	0
2	541	9	0	3	912	5	0	6	1101	3	0
6	1407	3	0	5	567	7	0	2	1207	3	0
2	725	7	0	8	868	4	0	4	1043	4	0
10	683	4	0	7	1308	3	0	2	685	6	0
1	683	7	0	7	1308	3	0	2	685	6	0
2	500	9	0	2	500	9	0	2	500	9	0
7	1308	3	0	1	1500	1	0	2	685	5	0
8	1500	2	0	1	1000	0	0	7	1000	3	0

Figure 4:  $A$ : packets sent in RX period,  $R$ : packets sent outside RX period

```
1 #!/bin/bash
2
3 for i in {0..9}
4 do
5     # Use one of the 10 predefined sink setups
6     ruby generate_sinks.rb $i
7
8     mrsh -i reset | grep "In_Reception" &
9
10    process=$!
11
12    # The extra 5 seconds allows leeway for setup etc.
13    sleep 65
14
15    killall mrsh
16    killall _mrsh
17 done
```